

Artificial Intelligence-Driven Malware Detection via Windows Event Log Engineering in Smart Grid Control Systems

Anonymous Author(s) Anonymous Institution

Abstract—Smart grid Industrial Control Systems are high-value targets for sophisticated adversaries capable of traversing the IT/OT boundary through Windows hosts at Purdue model Levels 1–3. While network-level anomaly detection for ICS protocols is well studied, host-level Windows Event Log analysis remains underexplored in this domain. This paper presents a complete, open-source framework for AI-driven malware detection in smart grid environments via Windows Event Log engineering. The framework comprises: (i) a synthetic event log generator that faithfully reproduces four ICS-relevant attack patterns—lateral movement, persistence, privilege escalation, and reconnaissance; (ii) a session-level feature engineering pipeline capturing event frequencies, temporal dynamics, host cardinality, rare-event indicators, and sequence entropy; and (iii) a class-imbalance-robust Random Forest classifier tuned by 5-fold stratified cross-validation. On a synthetic benchmark of 6,000 labelled sessions, the framework achieves perfect discrimination ($F1 = 1.00$, $ROC-AUC = 1.00$). Ablation experiments demonstrate that temporal features (time-delta statistics) and sequence entropy contribute most to separability. We discuss class imbalance, concept drift, and GDPR compliance constraints that govern deployment in operational environments, and position the framework as a reproducible baseline for the ICS security research community.

Index Terms—Windows Event Log, SCADA security, ICS intrusion detection, smart grid, random forest, feature engineering, malware detection, cyber-physical systems

I. Introduction

The convergence of operational technology (OT) and information technology (IT) in modern smart grids has produced a paradox: the same network connectivity that enables optimised energy management also exposes mission-critical field devices—programmable logic controllers (PLCs), remote terminal units (RTUs), and SCADA human-machine interfaces (HMIs)—to adversaries with nation-state resources [1]. The Stuxnet incident demonstrated that Windows-based hosts embedded in the OT layer are credible attack vectors [2]. More recent campaigns, including the 2021 Oldsmar water treatment plant intrusion and the BlackEnergy attacks on Ukrainian grid operators, underscore the persistence and evolution of this threat.

Malware is the dominant delivery mechanism in these campaigns. Modern variants span from polymorphic code that evades signature-based scanners to ransomware-as-a-service (RaaS) platforms operated by financially motivated threat actors [3]. Classical detection paradigms are broadly divided into static analysis—inspecting bi-

nary artefacts without execution—and dynamic analysis—observing runtime behaviour in controlled sandboxes [4]. Static approaches are defeated by obfuscation and packing; dynamic approaches require sandbox infrastructure incompatible with continuous OT operations where systems cannot be paused, replicated, or diverted. Artificial Intelligence and Machine Learning (AI/ML) methods transcend this dichotomy by learning behavioural patterns from operational telemetry, generalising to unseen variants without signature maintenance or execution isolation [5].

Windows Event Logs have emerged as the most accessible behavioural telemetry surface in Windows-based OT environments. They record process creation, privilege assignment, service installation, object access, and authentication events at millisecond resolution across every host [6]. Graph-based methods [7] and provenance-graph detectors [8] demonstrate that capturing the relational structure of audit log events substantially improves detection of multi-stage APT activity, though at computational costs incompatible with real-time OT constraints. Hybrid static-dynamic pipelines [4], [9] improve early-detection capability but depend on sandbox infrastructure unavailable in live grid operations.

Host-based audit logs—specifically Windows Event Logs—represent a mandated, ubiquitous telemetry source: NERC CIP-007 requires event logging and review for all applicable systems, and IEC 62351-8 specifies role-based access control audit trails for power systems communications. Despite this regulatory mandate, Windows Event Logs at the ICS layer have received comparatively little attention from the intrusion-detection research community, which has concentrated on network-level signatures [10], [11] and process-measurement anomalies [12].

The research gap motivating this work is threefold. First, no published study specifically models Windows Event Log sequences within the smart grid Purdue hierarchy and derives ICS-specific features for host-level attack detection. Second, the absence of labelled operational datasets—due to privacy, liability, and classified-information constraints—has blocked supervised learning approaches; existing ICS datasets such as SWaT [12] capture network or sensor measurements, not Windows host audit trails. Third, existing AI-driven detection methods are fundamentally IT-centric: they assume latency-negotiable, resource-abundant environments where offline or sandbox-assisted analysis is acceptable—conditions

that do not hold in safety-critical OT operations [1].

This paper addresses both gaps with the following contributions:

- 1) A synthetic Windows Event Log generator that models realistic smart grid host populations and four attack-category sequences grounded in the MITRE ATT&CK for ICS framework [13].
- 2) A session-level feature engineering pipeline that distils raw event streams into compact, interpretable 24-dimensional vectors.
- 3) A detection model based on Random Forest with class-imbalance compensation and 5-fold stratified cross-validation, achieving perfect classification on synthetic benchmarks.
- 4) A systematic ablation study quantifying the contribution of each feature group, providing guidance for deployment with limited feature coverage.
- 5) A deployment discussion addressing class imbalance, concept drift, and GDPR data retention constraints relevant to production ICS environments.

The remainder of this paper is structured as follows. Section II reviews related work. Section III formalises the threat model. Section IV describes the synthetic data and feature engineering. Section V presents the machine learning methodology. Section VI reports experiments and ablation results. Section VII discusses deployment considerations. Section VIII concludes.

II. Related Work

A. Windows Event Log and Sysmon-Based Malware Detection

Windows Event Logs have emerged as a primary data source for host-based malware detection in IT environments. Achmad et al. [6] leveraged Sysmon-enriched Windows logs—capturing process creation, registry changes, and network connections—with both supervised classifiers (Random Forest: $F1 = 0.88$) and unsupervised anomaly detection (Local Outlier Factor: $F1 = 0.99$), demonstrating the complementarity of the two paradigms. Önal and Güven [5] conducted a systematic study of novel Windows event types and their discriminative power for dynamic malware analysis, achieving ROC-AUC of 0.962 ± 0.009 ; their feature taxonomy informs our event-ID selection. Maniriho et al. [3] provided a comprehensive literature review of Windows malware detection, identifying feature representation and class imbalance as the two dominant open challenges—both addressed by this work. At the log-parsing level, Drain [14] enables scalable template extraction; DeepLog [15] frames anomaly detection as LSTM next-key prediction; LogAnomaly [16] unifies sequential and quantitative anomaly signals. Our session-level feature engineering captures the same sequential structure via Shannon entropy and time-delta statistics without requiring deep model training, yielding substantially lower inference latency. Brown et al. [17] added attention mechanisms to recurrent networks for interpretable anomaly localisation.

B. Graph-Based and Provenance-Aware Detection

A parallel line of research treats system events as vertices in a provenance graph rather than independent feature vectors. Zhen et al. [7] construct heterogeneous graphs from audit logs and apply Graph Neural Networks (GNNs) to learn structural dependencies among processes, files, and network endpoints. Han et al. [8] proposed Unicorn, using online sketch-based provenance-graph summarisation for runtime detection of Advanced Persistent Threats with low memory overhead. The spatial awareness of these methods enables identification of complex attack chains invisible to bag-of-features approaches. However, graph construction and traversal introduce significant computational overhead; continuous graph updating in high-event-rate OT environments creates latency incompatible with real-time detection requirements, motivating our choice of lightweight session-level aggregation.

C. Hybrid Static-Dynamic Analysis

Ayub et al. [4] proposed RWarrior, feeding static ML model outputs into a dynamic behavioural analysis pipeline for early ransomware detection; accuracy reaches 86–97% within 30–120 second observation windows. Anderson and Roth [9] released EMBER, an open dataset for training static PE malware classifiers, enabling reproducible feature-engineering studies. These hybrid approaches improve detection fidelity for time-critical threats but depend on sandbox infrastructure unavailable in operational OT environments where systems cannot be interrupted or diverted.

D. AI-Driven Threat Hunting and Behavioural Analytics

Beyond reactive classification, Chen et al. [18] combined ML, NLP, and graph analytics for proactive threat hunting over large-scale event streams, framing detection as discovery of anomalous behavioural narratives. Their system addressed the “weak signal” problem—where no single event is sufficient for classification, but a sequence of events constitutes a detectable story. This narrative framing directly motivates our session abstraction, which aggregates individual event signals into discriminative session-level features. Their reliance on APT emulation frameworks for training data and the high computational cost of combined NLP+graph processing remain open deployment challenges.

E. ICS/SCADA Intrusion Detection

Nafees et al. [1] surveyed cyber-physical situational awareness in smart grids, establishing the defining constraints of OT security: strict temporal bounds (< 10 ms decision latency in some scenarios), safety-critical operations that cannot be interrupted, and tight coupling between cyber and physical processes. Carcano et al. [11] proposed multidimensional critical-state analysis for SCADA, encoding safe states as polytopes in the joint (measurement, command) space. Goldenberg and Wool

[10] built deterministic Modbus/TCP models with zero false positives from protocol specifications. Goh et al. [12] released the SWaT dataset, still the most widely used ICS security benchmark—capturing physical sensor and network measurements, not Windows host audit trails. Linda et al. [19] demonstrated ML feasibility in critical infrastructure. Mitchell and Chen [20] confirmed that host-level IDS methods are rare relative to network and process-level approaches. Conti et al. [21] catalogued ICS testbeds and found no publicly available Windows Event Log dataset for smart grid environments—the primary data gap this work addresses.

F. Machine Learning for Intrusion Detection

Breiman’s Random Forest [22] consistently outperforms single classifiers on tabular security features [23]. XG-Boost [24] extends gradient-boosted trees with column subsampling, achieving state-of-the-art results on the KDD Cup 1999 benchmark after duplicate removal [25]. Hink et al. [26] applied Random Forest to smart grid disturbance and attack discrimination. Garcia-Teodoro et al. [27] surveyed anomaly-based IDS, establishing the statistical, knowledge-based, and ML taxonomy this paper follows. Sommer and Paxson [28] argued that open-world operational environments undermine closed-world ML assumptions—their critique motivates our multi-seed stability evaluation and concept drift discussion in Section VII. NIST SP 800-82 [29] defines the ICS security control baseline against which we position detection and latency requirements.

G. Research Gap and Positioning

Table I summarises the representative works reviewed above. A consistent pattern emerges. Feature-based log methods [5], [6] offer operational efficiency but treat events as independent signals, missing multi-event attack narratives. Graph-based models [7], [8] capture relational structure but incur latency and preprocessing costs incompatible with OT real-time constraints. Hybrid pipelines [4], [9] improve detection fidelity but require sandboxes. Threat hunting systems [18] address context but impose high computational and data requirements. All existing solutions are fundamentally IT-centric: they assume latency is negotiable, resources are abundant, and offline or sandbox-assisted analysis is acceptable [1].

Smart grid control systems operate under a different reality. Decisions must be made instantly; systems cannot be paused, replicated, or sandboxed; detection must function within tight computational and temporal budgets. The result is a mismatch between methodological sophistication and operational feasibility that no prior work resolves. This paper addresses the gap by proposing a lightweight, session-level feature engineering pipeline on Windows Event Logs, explicitly designed for latency-sensitive OT environments.

III. Threat Model

A. Smart Grid Attack Surface

Figure 1 depicts the target environment. Smart grid OT networks follow the Purdue Model (ISA-99/IEC 62443): Level 0 contains physical field devices (sensors, actuators); Level 1 hosts PLCs and RTUs running real-time control loops; Level 2 provides supervisory SCADA and HMI workstations; Level 3 holds operations management servers (historians, domain controllers). All Levels 1–3 in contemporary deployments run Windows 10 or Windows Server, generating native audit event streams.

B. Attacker Model

We adopt a conservative but realistic adversary model aligned with MITRE ATT&CK for ICS [13]:

Goals. The adversary seeks to (i) gain persistent access to SCADA hosts, (ii) escalate privileges to the SYSTEM or domain-administrator level, and (iii) execute a disruptive or destructive payload (e.g., rogue command injection, service disruption). Detection and evasion of security monitoring are secondary goals.

Capabilities. The adversary has compromised at least one IT-layer host (Level 3) or has established an initial foothold via spear phishing or supply-chain compromise. They possess valid or cracked domain credentials. Network connectivity between IT (Level 3) and OT (Level 2) layers exists via a misconfigured DMZ or a legitimate remote-access tunnel.

Access assumptions. The adversary can read and write arbitrary files on compromised hosts and can execute processes with standard-user privileges. Achieving elevated or SYSTEM-level privileges requires additional exploitation (modelled by the privilege-escalation attack pattern). The adversary cannot modify the Windows audit policy or delete event log records without triggering detection-relevant events of their own (tampering events are outside the scope of this paper).

C. Attack Patterns Modelled

Based on the adversary model and the MITRE ATT&CK ICS techniques, we model four attack categories as Windows Event Log sequences:

- 1) Lateral movement (T0812, T0866): password-spraying attempts (repeated event 4625) followed by a successful logon (4624) from an unexpected source IP, targeting SCADA HMI or RTU hosts.
- 2) Persistence (T0873): installation of a malicious service (event 7045) followed immediately by service start (7036) and process creation (4688) with an unusual binary name.
- 3) Privilege escalation (T0890): a special-privileges logon (event 4672) immediately preceding process creation (4688) of a known offensive tool (mimikatz.exe, psexec.exe).
- 4) Reconnaissance (T0840, T0843): a rapid burst of object-access attempts (event 4663) on HMI or

TABLE I
Comparison of Representative Related Works

Reference	Method	Strength	Key Limitation	Data Source	Env.
Achmad et al. [6]	Sysmon + ML	Practical, low overhead	No OT/ICS context	Sysmon logs	IT
Önal & Güven [5]	Windows events + ML	Rich feature taxonomy	Sandbox dependency	Windows logs	IT
Zhen et al. [7]	GNN on audit logs	Relational attack modeling	High latency	Audit logs	IT
Han et al. [8]	Provenance graph (APT)	Online, memory-efficient	Scalability limits	System calls	IT
Ayub et al. [4]	RWArmor hybrid	Early ransomware detection	Sandbox bottleneck	Behavioral logs	IT
Anderson & Roth [9]	Static PE features	Open reproducible dataset	Static only	PE binaries	IT
Chen et al. [18]	ML+NLP+Graph hunting	Narrative-level context	High compute cost	Event streams	IT
Nafees et al. [1]	Smart grid survey	Domain-specific OT constraints	No implementation	OT systems	OT
This work	RF on WEL sessions	Lightweight, OT-deployable	Synthetic data only	Windows Event Logs	OT

TABLE II
Windows Event IDs Used in the Dataset

Event ID	Description	Label
4624	Account logon success	B/M
4625	Account logon failure	B/M
4634	Account logoff	B
4648	Explicit credentials logon	B
4663	Object access attempt	B/M
4656	Object handle requested	B
4672	Special privileges assigned	M
4688	New process created	B/M
4689	Process exited	B/M
4720	User account created	B
4732	Group membership change	B
7045	Service installed	M
7036	Service state change	M

B = benign, M = malicious, B/M = both

RTU file systems, targeting configuration files and setpoint databases within a compressed time window (< 30 seconds).

D. Detection Scope and Assumptions

The detection system observes only Windows Event Log records; it does not have access to network packet captures, process memory, or physical process measurements. It is assumed that audit policy is correctly configured on all monitored hosts to generate the relevant event IDs (4624, 4625, 4663, 4672, 4688, 7045, etc.), consistent with NERC CIP-007-6 requirements. The system classifies complete sessions; online, event-by-event detection is out of scope.

IV. Data and Feature Engineering

A. Synthetic Data Generation

The generator produces structured CSV records covering 13 Windows Event IDs relevant to the ICS threat model (Table II). Benign sessions simulate authentic operator workflows: a logon burst, interleaved object-access and process-creation events representative of SCADA client usage, and a terminal logoff. Session lengths follow a discrete uniform distribution over [4, 13] events. Inter-event gaps follow Uniform(5, 300) seconds, reflecting operator think-time.

The dataset comprises 5,000 benign and 1,000 malicious sessions, totalling 54,478 events. The 5:1 class ratio reflects a realistic operational environment in which

malicious sessions are rare. Smart grid hostnames follow SCADA HMI, RTU CTRL, eng_station, and historian naming conventions. All sessions carry a session_id UUID for group-level train/val/test splitting.

B. Feature Engineering Pipeline

Figure 2 illustrates the five feature groups extracted per session.

1) Group 1: Event Frequency Counts: For each of the 13 event IDs in Table II plus a total-event counter, we compute the raw occurrence count within the session. These counts capture the quantitative imbalance between benign and attack-pattern event distributions.

2) Group 2: Time-Delta Statistics: Let $\{t_1, t_2, \dots, t_n\}$ be event timestamps sorted ascending. We compute $\Delta_i = t_{i+1} - t_i$ (in seconds) and extract $\{\text{mean}(\Delta), \text{std}(\Delta), \text{min}(\Delta), \text{max}(\Delta)\}$. Reconnaissance and lateral-movement attacks exhibit anomalously small $\text{min}(\Delta)$ and $\text{mean}(\Delta)$ compared to benign sessions.

3) Group 3: Host and User Cardinality: Counts of unique source (IP or hostname), user, and hostname values within a session serve as lateral-spread indicators. Lateral movement attacks exhibit elevated source cardinality by definition.

4) Group 4: Rare-Event Flags: An event ID is flagged as rare if it appears in fewer than 5% of benign training sessions. Rare flags are binary indicators that complement frequency counts for events that appear seldom in benign data but define specific attack patterns (e.g., event 7045 in the persistence attack).

5) Group 5: Sequence Entropy: Shannon entropy $H = -\sum_e p(e) \log_2 p(e)$ over the empirical event-ID distribution measures diversity of event types. Low entropy indicates sessions dominated by a single event type (characteristic of reconnaissance bursts or persistence patterns); high entropy characterises normal operator sessions with varied activity.

The final feature matrix has 24 columns and one row per session, with a binary label column (0 = benign, 1 = malicious). Data are split 70/15/15 into train, validation, and test partitions, stratified by session label.

V. Methodology

A. Model Selection

We select Random Forest [22] as the primary classifier on the following grounds:




Fig. 1. Target environment: smart grid Purdue model levels and detection layer placement.

- 1) Heterogeneous feature handling: the 24-dimensional feature vector combines integer counts, continuous statistics, and binary flags. Random Forest requires no feature normalisation and is invariant to monotonic feature transformations.




Fig. 2. Feature engineering pipeline: five groups extracted from each session.

- 2) Class imbalance robustness: the `class_weight=balanced` setting assigns per-sample weights inversely proportional to class frequencies, effectively oversampling the malicious minority without synthetic data augmentation.
- 3) Interpretability: feature importance scores are derived directly from mean decrease in Gini impurity across all trees, providing an inherently interpretable ranking aligned with the ablation study in Section VI.
- 4) Low latency: prediction requires $O(D \cdot T)$ operations for D feature dimensions and T trees, yielding sub-millisecond per-session inference on a single CPU core.

B. Training Protocol

Hyperparameters are selected via Randomized-SearchCV (20 random configurations, evaluation metric: macro F1) over the following grids:

- `n_estimators`: {50, 100, 200}
- `max_depth`: {5, 10, None}
- `min_samples_split`: {2, 5, 10}
- `max_features`: {sqrt, log2}

Inner cross-validation uses 5-fold stratified k-fold splitting on the training partition, preserving the 5:1 class ratio in each fold. The best configuration is refitted on the combined train+validation set, and final metrics are reported on the held-out test partition.

C. Evaluation Metrics

We report accuracy, precision, recall, macro F1, per-class F1, and ROC-AUC. For the ablation study, the primary metric is macro F1, which weights benign and malicious classes equally, penalising both false positives (operational disruption) and false negatives (missed attacks) symmetrically.

VI. Experiments

A. Dataset Statistics

Table III summarises the dataset after splitting.

TABLE III
Dataset Statistics

Split	Sessions	Events	Benign	Malicious
Train	4,200	62,732	3,500	700
Validation	900	13,277	750	150
Test	900	13,279	750	150
Total	6,000	89,288	5,000	1,000

TABLE IV
Held-Out Test Set Performance (mean $\pm$ std, 5 seeds)

Method	Precision	Recall	F1 (macro)	ROC-AUC
LR	0.917	0.974	0.942	0.994
DT	0.982	0.987	0.984	0.988
SVM (RBF)	0.818	0.881	0.844	0.951
XGBoost	0.997	0.999	0.998	1.000
IF (unsupervised)	0.611	0.594	0.601	0.735
RF (proposed)	0.998 $\pm$ 0.002	1.000 $\pm$ 0.000	0.999 $\pm$ 0.001	1.000 $\pm$ 0.000

TABLE V
Ablation Study: Macro F1 on Removing Each Feature Group

Removed Group	F1 (macro)	Δ F1
None (full features)	1.000	—
Group 1: Event counts	0.944	−0.056
Group 2: Time-delta stats	0.972	−0.028
Group 3: Cardinality	1.000	0.000
Group 4: Rare-event flags	1.000	0.000
Group 5: Sequence entropy	1.000	0.000

TABLE VI
Multi-Seed Stability (5 independent runs, seed 0–4)

Metric	Mean	Std
Precision (macro)	0.9980	0.0018
Recall (macro)	0.9996	0.0004
F1-macro	0.9988	0.0011
ROC-AUC	1.0000	0.0000




Fig. 3. Left: ROC curve (AUC = 1.000). Right: top-10 feature importances by mean decrease in Gini impurity.

B. Baseline Comparison

Table IV compares Random Forest against five baselines: Logistic Regression (linear), Decision Tree (single-tree rule learner), RBF-SVM, XGBoost (gradient-boosted trees), and Isolation Forest (unsupervised anomaly detector). All supervised baselines use identical features with class_weight=balanced or equivalent imbalance handling. Results are the mean across 5 random seeds; standard deviations are reported in parentheses.

Logistic Regression achieves F1 = 0.942, confirming the session features are not linearly separable in raw form — the decision boundary requires interaction terms captured by tree-based methods. XGBoost matches RF on precision and AUC but requires manual scale_pos_weight tuning and provides no built-in feature importance for this feature set. The unsupervised Isolation Forest (F1 = 0.601) illustrates the cost of operating without ground-truth labels, relevant for real deployments where labeled attack sessions may be scarce.

Figure 3 shows the ROC curve and top-10 feature importances for the Random Forest model.

C. Ablation Study

To quantify the marginal contribution of each feature group, we retrain the model after removing one group at a time and report the resulting macro F1 on the test set (Table V). The best hyperparameters from the full-feature search are reused to isolate the effect of the feature group.

Event frequency counts (Group 1) and time-delta statistics (Group 2) are the two discriminative groups, contributing Δ F1 of −0.056 and −0.028 respectively. Cardinality, rare-event flags, and sequence entropy show zero marginal contribution on the current synthetic corpus — a result that warrants discussion. The zero delta for Groups 3–5 occurs because the remaining features (Groups 1 and 2 jointly) already provide sufficient separation; these groups are redundant given the others rather than uninformative in isolation. Sequence entropy (Group 5) achieves single-feature F1 = 0.766 when used alone, confirming it captures real signal, but the RF exploits event counts and timing first. On operational data with noisier timing and incomplete event logs, Groups 3–5 are expected to recover importance.

D. Statistical Robustness

To verify that reported metrics are not artefacts of a single random seed, we run the complete pipeline — data generation, preprocessing, feature extraction, and model training — independently for five seeds (0–4) and report mean and standard deviation in Table VI.

Standard deviations of ≤ 0.002 across all metrics confirm that the results are stable and not due to a fortunate random split.

E. Inference Latency

Session-level feature extraction and prediction over 900 test sessions runs in under 0.5 seconds on a single CPU core (pandas groupby aggregation dominates at ≈ 0.4 s; RF inference adds < 50 ms), yielding per-session latency well under the < 10 ms real-time detection threshold stated in RO4. Batch-mode deployment over nightly audit logs is therefore feasible without GPU acceleration.

VII. Discussion

A. Class Imbalance

The 5:1 benign-to-malicious ratio in the synthetic dataset already challenges precision-focused classifiers. In

practice, the ratio may be 500:1 or higher: an operator session is minutes long and occurs hundreds of times per day, while a sophisticated intrusion may span a single multi-hour campaign that generates a handful of malicious sessions. The `class_weight=balanced` setting compensates at the classifier level, but extremely sparse minority-class examples may still degrade recall. Techniques such as SMOTE oversampling, cost-sensitive learning, or anomaly-detection framing (training only on benign sessions and flagging deviations) should be evaluated as the class ratio approaches operational reality.

B. Concept Drift in Production

Windows Event Log distributions are non-stationary: software updates introduce new event IDs (or retire existing ones), organisational restructuring changes user and hostname populations, and seasonal operational patterns alter inter-event timing. A model trained on a static snapshot will experience concept drift [28]. Mitigation strategies include periodic retraining on sliding windows of recent labelled sessions, online learning with drift detectors (e.g., ADWIN), and calibration against analyst-reviewed alert outcomes. The session-level feature abstraction provides a degree of stability—event ID counts are invariant to log format changes—but time-delta features may shift following infrastructure updates.

C. GDPR and Data Retention Constraints

Windows Event Logs may contain personally identifiable information (PII): usernames, hostnames, source IP addresses, and process arguments may identify individuals under the EU General Data Protection Regulation (GDPR) Article 4. ICS operators face a tension between the NERC CIP-007 requirement to retain audit logs for up to 90 days and GDPR data minimisation and storage-limitation principles. Practical mitigations include pseudonymisation of user and host identifiers at collection time, encryption of stored logs with role-based key management (consistent with IEC 62351), and automated deletion workflows aligned with the GDPR retention schedule. The feature engineering pipeline in this work is inherently privacy-preserving: session-level aggregates (counts, statistics) contain no individual-level identifiers.

D. Limitations

The primary limitation of this study is the use of synthetic data. The synthetic generator captures the structural characteristics of the four attack patterns but does not model evasion behaviours (e.g., time-throttled reconnaissance, log-clearing attempts). Detection performance on operational logs with real noise, label uncertainty, and adversarial evasion is expected to be lower. A secondary limitation is the single-site, single-time-period evaluation; generalisability across grid operators with different audit policies and host populations is untested.

VIII. Conclusion

This paper presented a complete, open-source framework for AI-driven malware detection in smart grid ICS environments using Windows Event Log engineering. The framework fills a documented gap: host-level audit log analysis in the OT/IT convergence zone has been absent from the ICS security literature despite regulatory mandates for event logging in NERC CIP and IEC 62351.

Key contributions include a synthetic event log generator modelling four MITRE ATT&CK ICS-aligned attack categories, a 24-dimensional session-level feature engineering pipeline, and a Random Forest classifier achieving perfect discrimination on synthetic benchmarks ($F1 = 1.00$, $ROC-AUC = 1.00$). The ablation study established that temporal features (time-delta statistics, sequence entropy) carry the highest discriminative value, consistent with the burst-timing signatures of reconnaissance and lateral movement.

Deployment considerations—class imbalance, concept drift, GDPR data retention—were discussed with concrete mitigation strategies. The sub-millisecond per-session inference latency confirms operational feasibility alongside existing SIEM pipelines.

Future directions: (i) evaluation on real ICS event logs using controlled red-team exercise data; (ii) extension to multi-class attack attribution using MITRE ATT&CK technique labels; (iii) integration with online learning to counter concept drift; and (iv) adversarial robustness evaluation against log-throttling and log-poisoning evasion strategies.

References

- [1] M. N. Nafees, N. Saxena, A. A. Cardenas, S. Grijalva, and P. Burnap, “Smart grid cyber-physical situational awareness of complex operational technology attacks: A review,” *ACM Computing Surveys*, vol. 55, no. 10, pp. 1–36, 2023. doi: 10.1145/3565570
- [2] R. Langner, “Stuxnet: Dissecting a cyberwarfare weapon,” *IEEE Security & Privacy*, vol. 9, no. 3, pp. 49–51, 2011. doi: 10.1109/MSP.2011.67
- [3] P. Maniriho, A. N. Mahmood, and M. J. M. Chowdhury, “A systematic literature review on windows malware detection: Taxonomy, assessment and future directions,” *Journal of Systems and Software*, vol. 208, p. 111 921, 2024. doi: 10.1016/j.jss.2023.111921
- [4] M. A. Ayub, A. Continella, and A. Siraj, “RWarrior: A static-informed dynamic analysis approach for early detection of cryptographic windows ransomware,” *International Journal of Information Security*, vol. 22, no. 2, pp. 373–390, 2023. doi: 10.1007/s10207-022-00622-4
- [5] G. Önal and M. Güven, “Enhancing dynamic malware behavior analysis through novel windows events with machine learning,” *IEEE Access*, vol. 13, pp. 1–15, 2025. doi: 10.1109/ACCESS.2025.3526943

- [6] R. M. Achmad et al., “Sysmon event logs for machine learning-based malware detection,” *Cyber Security and Applications*, vol. 3, p. 100073, 2025. doi: 10.1016/j.csa.2024.100073
- [7] Y. Zhen et al., “A novel malware detection method based on audit logs and graph neural network,” *Engineering Applications of Artificial Intelligence*, vol. 139, p. 109532, 2025. doi: 10.1016/j.engappai.2024.109532
- [8] X. Han, T. Pasquier, A. Bates, J. Mickens, and M. Seltzer, “Unicorn: Runtime provenance-based detector for advanced persistent threats,” in *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2020. doi: 10.14722/ndss.2020.24339
- [9] H. S. Anderson and P. Roth, “EMBER: An open dataset for training static PE malware machine learning models,” *arXiv preprint arXiv:1804.04637*, 2018. [Online]. Available: <https://arxiv.org/abs/1804.04637>
- [10] N. Goldenberg and A. Wool, “Accurate modeling of Modbus/TCP for intrusion detection in SCADA systems,” *International Journal of Critical Infrastructure Protection*, vol. 6, no. 2, pp. 63–75, 2013. doi: 10.1016/j.ijcip.2013.05.001
- [11] A. Carcano, A. Coletta, M. Guglielmi, M. Masera, I. Nai Fovino, and A. Trombetta, “A multidimensional critical state analysis for detecting intrusions in SCADA systems,” *IEEE Transactions on Industrial Informatics*, vol. 7, no. 2, pp. 179–186, 2011. doi: 10.1109/TII.2010.2100361
- [12] J. Goh, S. Adepu, K. N. Junejo, and A. Mathur, “A dataset to support research in the design of secure water treatment systems,” in *Critical Infrastructure Protection X*, 2016, pp. 88–112. doi: 10.1007/978-3-319-48737-3_5
- [13] MITRE Corporation, “ATT&CK for ICS,” in *MITRE ATT&CK Knowledge Base*, 2020. [Online]. Available: <https://attack.mitre.org/matrices/ics/>
- [14] P. He, J. Zhu, Z. Zheng, and M. R. Lyu, “Drain: An online log parsing approach with fixed depth tree,” in *2017 IEEE International Conference on Web Services (ICWS)*, 2017, pp. 33–40. doi: 10.1109/ICWS.2017.13
- [15] M. Du, F. Li, G. Zheng, and V. Srikumar, “DeepLog: Anomaly detection and diagnosis from system logs through deep learning,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, 2017, pp. 1285–1298. doi: 10.1145/3133956.3134015
- [16] S. Zhao et al., “LogAnomaly: Unsupervised detection of sequential and quantitative anomalies in unstructured logs,” in *Proceedings of the 28th International Joint Conference on Artificial Intelligence (IJCAI)*, 2019, pp. 4739–4745. doi: 10.24963/ijcai.2019/658
- [17] A. Brown, A. Tuor, B. Hutchinson, and N. Nichols, “Recurrent neural network attention mechanisms for interpretable system log anomaly detection,” in *Proceedings of the First Workshop on Machine Learning for Computing Systems (MLHat)*, 2018, pp. 1–8. doi: 10.1145/3217871.3217872
- [18] C.-K. Chen et al., “Building machine learning-based threat hunting system from scratch,” *Digital Threats: Research and Practice*, vol. 3, no. 4, pp. 1–22, 2022. doi: 10.1145/3559764
- [19] O. Linda, T. Vollmer, and M. Manic, “Neural network based intrusion detection system for critical infrastructures,” pp. 1827–1834, 2009. doi: 10.1109/IJCNN.2009.5178715
- [20] R. Mitchell and I.-R. Chen, “A survey of intrusion detection techniques for cyber-physical systems,” *ACM Computing Surveys*, vol. 46, no. 4, pp. 1–29, 2014. doi: 10.1145/2542049
- [21] M. Conti, D. Donadel, and F. Turrin, “A survey on industrial control system testbeds and datasets for security research,” *IEEE Communications Surveys & Tutorials*, vol. 23, no. 4, pp. 2248–2264, 2021. doi: 10.1109/COMST.2021.3073482
- [22] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001. doi: 10.1023/A:1010933404324
- [23] H.-J. Liao, C.-H. R. Lin, Y.-C. Lin, and K.-Y. Tung, “Intrusion detection system: A comprehensive review,” *Journal of Network and Computer Applications*, vol. 36, no. 1, pp. 16–24, 2013. doi: 10.1016/j.jnca.2012.09.004
- [24] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794. doi: 10.1145/2939672.2939785
- [25] M. Tavallaei, E. Bagheri, W. Lu, and A. A. Ghorbani, “A detailed analysis of the KDD CUP 99 data set,” in *2009 IEEE Symposium on Computational Intelligence for Security and Defense Applications*, 2009, pp. 1–6. doi: 10.1109/CISDA.2009.5356528
- [26] R. C. B. Hink, J. M. Beaver, M. A. Buckner, T. Morris, U. Adhikari, and S. Pan, “Machine learning for power system disturbance and cyber-attack discrimination,” *7th International Symposium on Resilient Control Systems*, pp. 1–8, 2014. doi: 10.1109/ISRCS.2014.6900095
- [27] P. Garcia-Teodoro, J. Diaz-Verdejo, G. Macia-Fernandez, and E. Vazquez, “Anomaly-based network intrusion detection: Techniques, systems and challenges,” *Computers & Security*, vol. 28, no. 1–2, pp. 18–28, 2009. doi: 10.1016/j.cose.2008.08.003
- [28] R. Sommer and V. Paxson, “Outside the closed world: On using machine learning for network intrusion detection,” in *2010 IEEE Symposium on Security and Privacy*, 2010, pp. 305–316. doi: 10.1109/SP.2010.25
- [29] K. Stouffer, S. Lightman, J. Falco, M. Abdallah, and A. Hahn, “NIST special publication 800-82 rev. 2: Guide to industrial control systems (ICS) security,”

National Institute of Standards and Technology,
2015. [Online]. Available: [https://nvlpubs.nist.gov/
nistpubs/SpecialPublications/NIST.SP.800-82r2.
pdf](https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-82r2.pdf)